

Inside the Brain

Build a neural network by hand — nothing but NumPy

Wednesday · July 22, 2026

Engine day · no module — you forge the core


1 Mission briefing

The Chief Scientist drops by with a grin: "Anyone can call a library. Today you build a neural network with your bare hands — just NumPy — and train it to think." No PyTorch, no shortcuts. *Previously:* you unlocked the Pattern Finder and watched machines find structure with no labels. Today you build the thing everyone means when they say "AI" — and after this, the black box is never black again.

- ☐ Understand a **neuron** and why one is never single enough
- ☐ Discover a layer is **one matrix multiplication** really
- ☐ Write **forward**, **loss**, and **nudge** update the **pass** a the
- ☐ *Watch* the decision boundary bend around the data as it learns
- ☐ See **hidden** and **learning** change how **size** **rate** everything
- ☐ Peek at **loss** the network rolls the **landscape** down

2 Field guide the concepts

A neuron is a bouncer at a door. Each input (a reason to open) has a **weight** (how much it matters). The bouncer multiplies each input by its weight, sums them, adds a personal **bias**, and runs the total through an **activation** — the **sigmoid** squash that turns any number into a 0-to-1 confidence. One neuron draws one straight cut, so it can't split two tangled moons; stack a **hidden layer** of them and the boundary can bend.



One neuron. A whole **layer** of them is a single matrix multiply: $(300, 2) @ (2, 16) \rightarrow (300, 16)$.

Learning = walking downhill. A **loss** scores how wrong the guesses are (one number). The **gradient** says which way each weight should move to shrink it; today you estimate it by *nudging* each weight and re-checking the loss, then step every weight a little downhill (the **learning rate** sets the step size). Remember the big secret: **everything the network does is matrix multiplication** — Weeks 3–4 make those fast.

3 Lab missions check each off in the notebook

1 Implement the sigmoid

`1 / (1 + np.exp(-z))`, working on a whole array at once.

Expect: `sigmoid(0) = 0.5` and a smooth S-curve from 0 to 1.



2 Matrix multiply, by hand vs NumPy



Write the triple-loop `matmul_by_hand`, then race it against `@`.

Expect: the exact same answer — and NumPy winning by hundreds of times.

3 Write the forward pass



Chain two layers: `H = sigmoid(X@W1+b1)`, then the output layer.

Expect: a (300, 1) array of mush near 0.5 — it hasn't learned yet.

4 Implement the loss



Mean squared error: the average of `(pred - truth)2`.

Expect: an untrained loss near 0.25.

5 Take the step (this is learning)



Nudge every weight downhill, in place: `p -= lr * g`.

Expect: the loss falls epoch by epoch as the net figures out the moons.

6 Score it



Threshold predictions at 0.5 and take the fraction correct.

Expect: a well-trained net lands around 88% on the moons (the ceiling is ~97% — tomorrow's power tools close the gap).

7 Retrain with 2 and 64 hidden neurons



Train both and plot their boundaries side by side.

Expect: 2 is too stiff to bend; 64 is far wigglier.

4 Cheat sheet a neural net in pure NumPy

<code>sigmoid(z) = 1 / (1 + np.exp(-z))</code>	Squash any number to a 0–1 confidence
<code>X @ W1 + b1</code>	A whole layer of neurons in one matrix multiply
<code>(300,2) @ (2,16) → (300,16)</code>	Shapes: data @ weights → hidden values
<code>sigmoid(sigmoid(X@W1+b1) @ W2 + b2)</code>	The full two-layer forward pass
<code>np.mean((y_pred - y_true) ** 2)</code>	MSE loss — one number of wrongness
<code>p -= lr * g</code>	The update: step each weight downhill (learning!)
<code>loss(w+eps) vs loss(w-eps)</code>	The nudge: estimate a weight's gradient
<code>np.random.default_rng(SEED)</code>	Reproducible small random starting weights

5 What you forged today

The engine — built by hand

No file to save today. You built the **core** that tomorrow's Studio module runs on: a neural network from nothing but NumPy that teaches itself to separate two interlocking moons.

65

weights & biases
($2 \cdot 16 + 16 + 16 + 1$)

~88%

accuracy on the
two-moons task

130 → 2

forward passes per step:
nudge today, autograd tomorrow

Tomorrow the Lab bolts on power tools — PyTorch autograd and a GPU — and this same engine learns to read handwriting.

6 Mini-glossary

neuron

Weighted sum of inputs + bias, run through an activation.

bias

A neuron's offset — its baseline mood, added before the squash.

hidden layer

Neurons between input and output; more = a wigglier boundary.

gradient

The slope of the loss — which way to move each weight to improve.

weight

How much one input counts — a knob the network tunes.

activation

The squashing step (sigmoid) that bends straight lines into curves.

loss

One number saying how wrong the predictions are (here, MSE).

learning rate

How big a step to take downhill each update.

7 Tonight's show & tell

1. How many times did NumPy beat your hand-written **matmul**? Why does that gap decide the back half of this course?
2. Show the boundary at **2 vs 64** hidden neurons — which one memorized, which generalized?
3. The nudge method spent ~**130 forward passes** per step. What does tomorrow's **autograd** promise instead?